

Working with Unlabeled Data – Cluster Analysis

Find the best number of clusters with **k_means** and **agglomerative clustering**

Overview

1. Load the data file
 - check the shape and plot the content
2. Observe the pair plot and comment the shapes in view of clustering
 - A. if necessary, transform the data
3. Use the elbow method to find the optimal number of clusters, to do this test `KMeans` with varying number of clusters, from 2 to 10: for each value of `k`
 - fit the data
 - compute the **inertia** and the **silhouette score**
 - store them for plot
4. Plot inertia and silhouette score versus `k`
5. Choose the optimal number of clusters looking at the plots
6. Cluster the data using the optimal number, plot the cluster assignment
 - in the plot choose the features that seem to be most promising

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score, silhouette_samples
from mpl_toolkits import mplot3d
from sklearn.model_selection import ParameterGrid
import warnings
warnings.filterwarnings("ignore")

random_state = 42 # This variable will be used in all the procedure calls allowing
                  # in this way the running can be perfectly reproduced
                  # just change this value for a different experiment
```

1. Load the data file

Check the shape and plot the content

```
In [2]: X_url = 'https://archive.ics.uci.edu/ml/machine-learning-databases/00292/Wholesale'
delimiter = ','
X0 = pd.read_csv(X_url, delimiter=delimiter)
X0.shape
```

Out[2]: (440, 8)


```
In [3]: X0.head()
```

Out[3]:


	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
0	2	3	12669	9656	7561	214	2674	1338
1	2	3	7057	9810	9568	1762	3293	1776
2	2	3	6353	8808	7684	2405	3516	7844
3	1	3	13265	1196	4221	6404	507	1788
4	2	3	22615	5410	7198	3915	1777	5185

2. Observe the data distributions

```
In [4]: X0.boxplot(figsize=(15,10));
```



```
In [5]: sns.pairplot(X0);  
plt.show()
```



We observe that the distributions of values are definitely *skewed*: in the columns from `Fresh` to `Delicassen` the values are highly concentrated on the right, but there are always outliers, frequently in a very large range.

Clustering is more effective in absence of outliers and with all the variables distributed in similar ranges, for this reason, we will execute two transformations:

1. transform all the variables from the column `Fresh` to the column `Delicassen` with `PowerTransformer`
2. remap all the variables in the range `0:1`

```
In [6]: from sklearn.preprocessing import PowerTransformer
pt = PowerTransformer()
X_pt = pd.DataFrame(pt.fit_transform(X0.iloc[:,2:]),columns=X0.columns[2:])
X_trasf = pd.concat([X0.iloc[:,2:],X_pt],axis=1)
from sklearn.preprocessing import MinMaxScaler
mms = MinMaxScaler()
X = pd.DataFrame(mms.fit_transform(X_trasf), columns = X_trasf.columns)
X.head()
```

Out[6]:


	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
0	1.0	1.0	0.501828	0.667606	0.625238	0.208640	0.649941	0.4120
1	1.0	1.0	0.414266	0.670028	0.655690	0.458800	0.674852	0.4456
2	1.0	1.0	0.400077	0.653586	0.627297	0.499856	0.682752	0.6542
3	0.0	1.0	0.509368	0.369264	0.553550	0.636716	0.461095	0.4464
4	1.0	1.0	0.604755	0.580657	0.618985	0.566470	0.601884	0.5902



Show the result of the transformation


In [7]:

```
X.boxplot(figsize=(15,8));
plt.show()
```



In [8]:

```
sns.pairplot(X);
plt.show()
```



Now the effect of outliers is reduced, and we compute the clustering

3. Use the elbow method to find the optimal number of clusters

Test `KMeans` with varying number of clusters, from 2 to 10

Prepare the results list that will contain pairs of `inertia` and `silhouette_score` for each value of `k`, then, **for each value of `k`**

- initialize an estimator for `KMeans`
- fit the data and predict the cluster assignment for each individual with `fit` and `predict`
- the **inertia** is provided in the attribute `inertia_` of the fitted model
- compute the **silhouette score** using the function `silhouette_score` from `sklearn.metrics` using as arguments the data and the fitted labels, we will fill the variable `silhouette_scores`
- store the two values above in the list created at the beginning

```
In [9]: k_range = list(range(2,11)) # set the range of k values to test
parameters_km = [{'n_clusters': k_range}]
pg = list(ParameterGrid(parameters_km))
inertias_km = []
silhouette_scores_km = []
for i in range(len(pg)):
    km = KMeans(**(pg[i]), random_state=random_state)
    y_km = km.fit_predict(X)
    inertias_km.append(km.inertia_)
    silhouette_scores_km.append(silhouette_score(X,y_km))
```

4. Plot inertia and silhouette score versus k


```
In [10]: def two_plots(x, y1, y2, xlabel, y1label, y2label):
    fig, ax1 = plt.subplots()

    color = 'tab:red'
    ax1.set_xlabel(xlabel)
    ax1.set_ylabel(y1label, color=color)
    ax1.plot(x, y1, color=color)
    ax1.tick_params(axis='y', labelcolor=color)
    ax2 = ax1.twinx() # instantiate a second axes that shares the same x-axis

    color = 'tab:blue'
    ax2.set_ylabel(y2label, color=color) # we already handled the x-Label with
    ax2.plot(x, y2, color=color)
    ax2.tick_params(axis='y', labelcolor=color)
    ax2.set_ylim(0,1) # the axis for silhouette is [0,1]

    fig.tight_layout() # otherwise the right y-label is slightly clipped
    plt.show()
```

```
In [11]: two_plots(x=k_range, y1=inertias_km, y2=silhouette_scores_km
                 , xlabel='Number of clusters', y1label='Inertias', y2label='Silhouette')
```



5. Cluster with the optimal number

The two *elbow* points of inertia would suggest as cluster number 3 or 4, slightly more pronounced in 3. Silhouette has a maximum on 4, but the increase with respect to 3 is very small.


We will choose $k=4$

```
In [12]: k=4
km = KMeans(n_clusters=k,
             random_state=random_state)
y_km = km.fit_predict(X)
print("Number of clusters = {}\t- Distortion = {:.2f}\t- Silhouette score = {:.4f}"
      .format(k,inertias_km[k_range.index(k)],silhouette_scores_km[k_range.index(k)])
```

Number of clusters = 4 - Distortion = 54.49 - Silhouette score = 0.49

Show the distribution of samples in the clusters with a pie chart

```
In [13]: clust_sizes_km = np.unique(y_km,return_counts=True)
pd.DataFrame(clust_sizes_km[1]).plot.pie(y=0, autopct='%1.1f%%', );
plt.show()
```



Comments

The **silhouette score** ranges from `-1` (worst) to `1` (best); as a rule of thumb, a value greater than `0.5` should be considered acceptable.